

# Introduction to BigData Computing Framework - Apache Hadoop & Spark


國網中心  
莊家雋

# What will you learn today ?

- Distributed Computing environment
- Apache Hadoop
  - Concept of MapReduce
  - MapReduce Programing
  - MapReduce Examples
- Apache Spark
  - Concept of RDD
  - RDD Programing
  - Spark Examples

# Distributed System

- Master /slave architecture




- 開發一個分散式系統很難
  - 主機間如何溝通
  - 系統的可靠性設計
  - ...
- Hadoop
  - HDFS、MapReduce
- Spark
  - RDD




# 分散式檔案系統：HDFS

- 在分散式的儲存環境裏，提供邏輯上的單一目錄系統
- 每個檔案被分割成許多區塊並進行異地備份

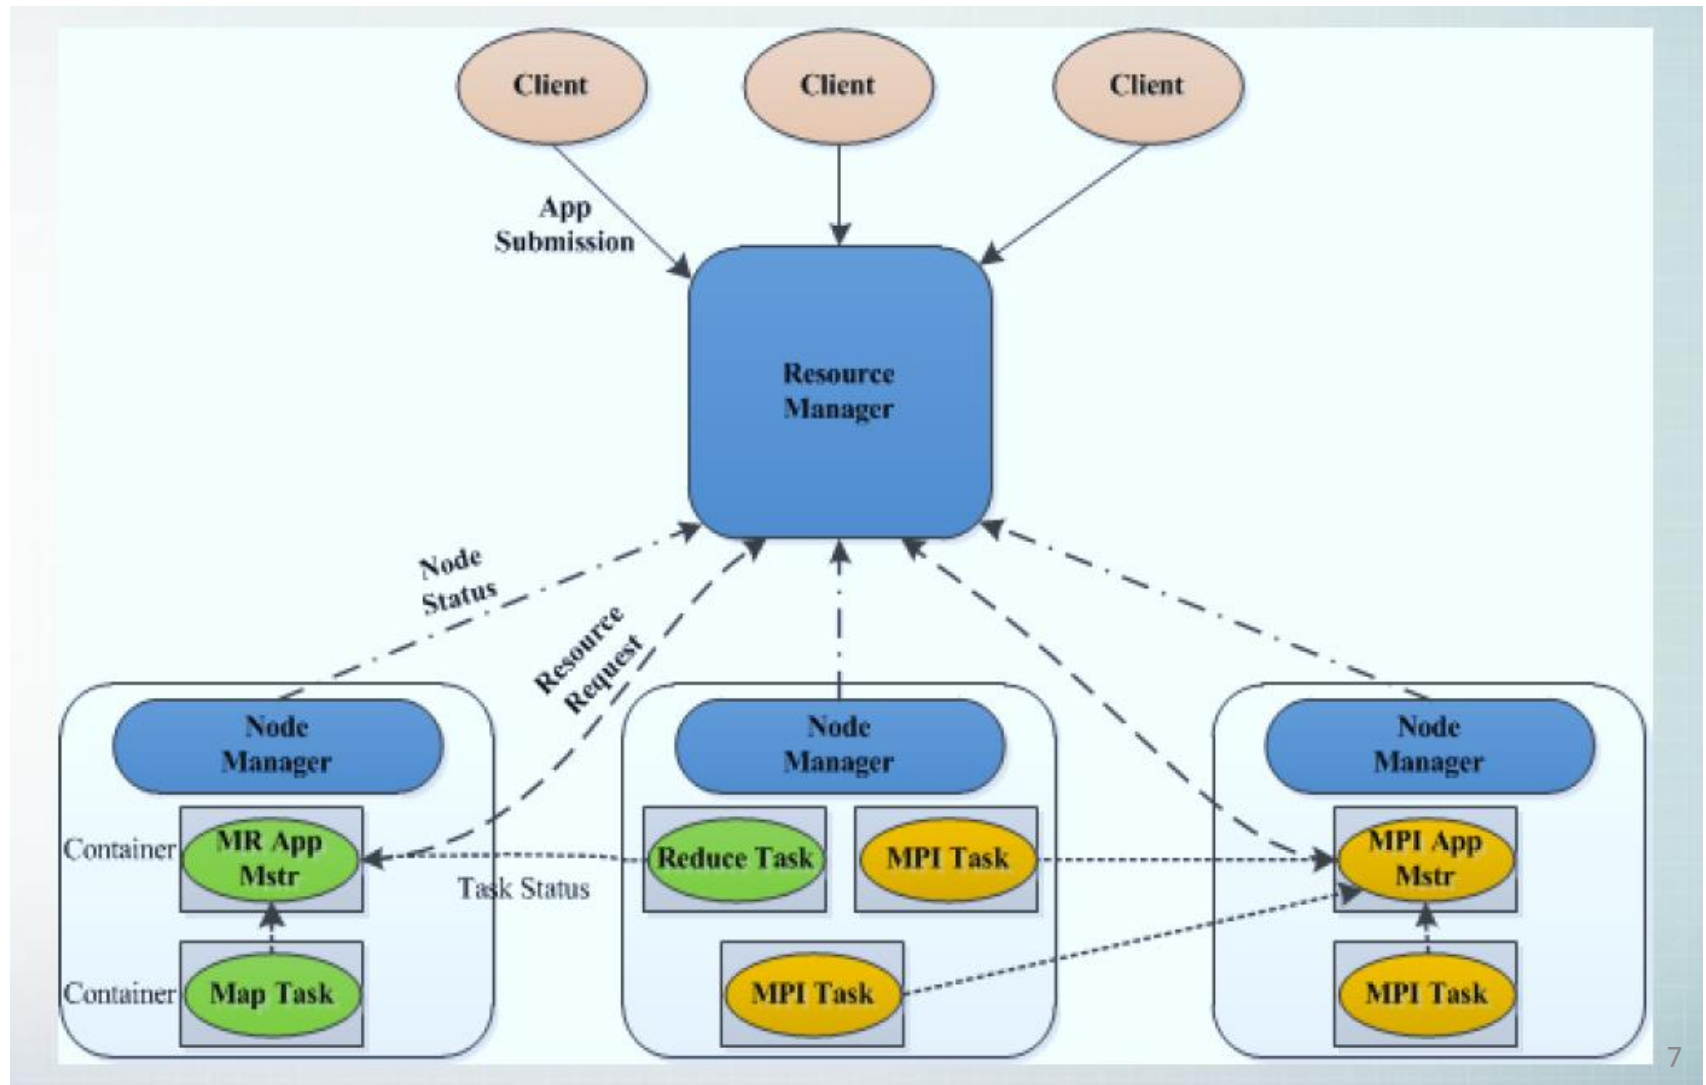


# Data Locality

- 移動運算到資料端比移動資料到運算端來的成本低
  - 減少資料搬運，實現在地運算



# Concept of Distributed computing



# MapReduce Concept & Real life Example







# 數學歸納法證明

步驟 1、證明 $n=1$  時，敘述成立。

步驟 2、假設 $n=k$ 時，敘述成立；證明 $n=k+1$ ，敘述也成立。

由數學歸納法得證， $n$  為任意自然數時都成立

3. Show that any  $2^n \times 2^n$  board with one square deleted can be covered by Triominoes.



A Triominoe

試證明：當自然數  $n \geq 3$  時，不等式： $5^n > 3^n + 4^n$  恆成立.

證明：(1)  $n=3$  時，左式  $= 5^3 = 125$ ，右式  $= 3^3 + 4^3 = 91$ ，因  $125 > 91$ ，  
故  $5^3 > 3^3 + 4^3$  成立.

(2) 設  $n=k$  ( $k$  是一個整數且  $k \geq 3$ ) 時， $5^k > 3^k + 4^k$  成立.

上式兩端同乘 5 得：

$$5^{k+1} > 5 \cdot 3^k + 5 \cdot 4^k > 3 \cdot 3^k + 4 \cdot 4^k = 3^{k+1} + 4^{k+1},$$

故  $5^{k+1} > 3^{k+1} + 4^{k+1}$  亦成立.

由數學歸納法知  $5^n > 3^n + 4^n$  ( $n \geq 3$ ) 成立.


# 選舉到了...

- 台北市10個選區，共100萬票，要算出每個候選人的得票數




監票人1  
[負責1區]

號次	票數
2	1
1	1
...	...




監票人2  
[負責2區]

號次	票數
1	1
1	1
3	1
...	...




監票人3  
[負責3區]

號次	票數
3	1
2	1
1	1



監票人4  
[負責4區]

號次	票數
1	1
3	1
3	...



監票人5  
[負責5區]

號次	票數
3	1
2	1
3	1

號次	票數
2	1
1	1
...	...

號次	票數
5	1
1	1
7	1
...	...

號次	票數
5	1
2	1
1	1


號次	票數
1	1
5	1
3	...

號次	票數
4	1
2	1
6	1

Shuffle & Sort  
由各投開票所送到中選會

號次	票數	號次	票數	號次	票數
1	1	2	1	3	1
1	1	2	1	3	1
1	1	2	1	3	1
1	1	2	1	3	1
1	...	2	...	3	...

中選會  
[負責全部的候選人]



號次	總票數
1	187532

號次	總票數
2	574821

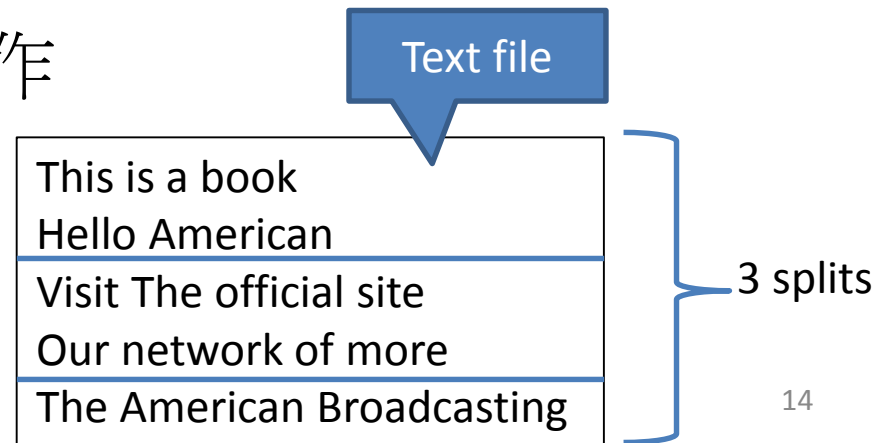
號次	總票數
3	237647

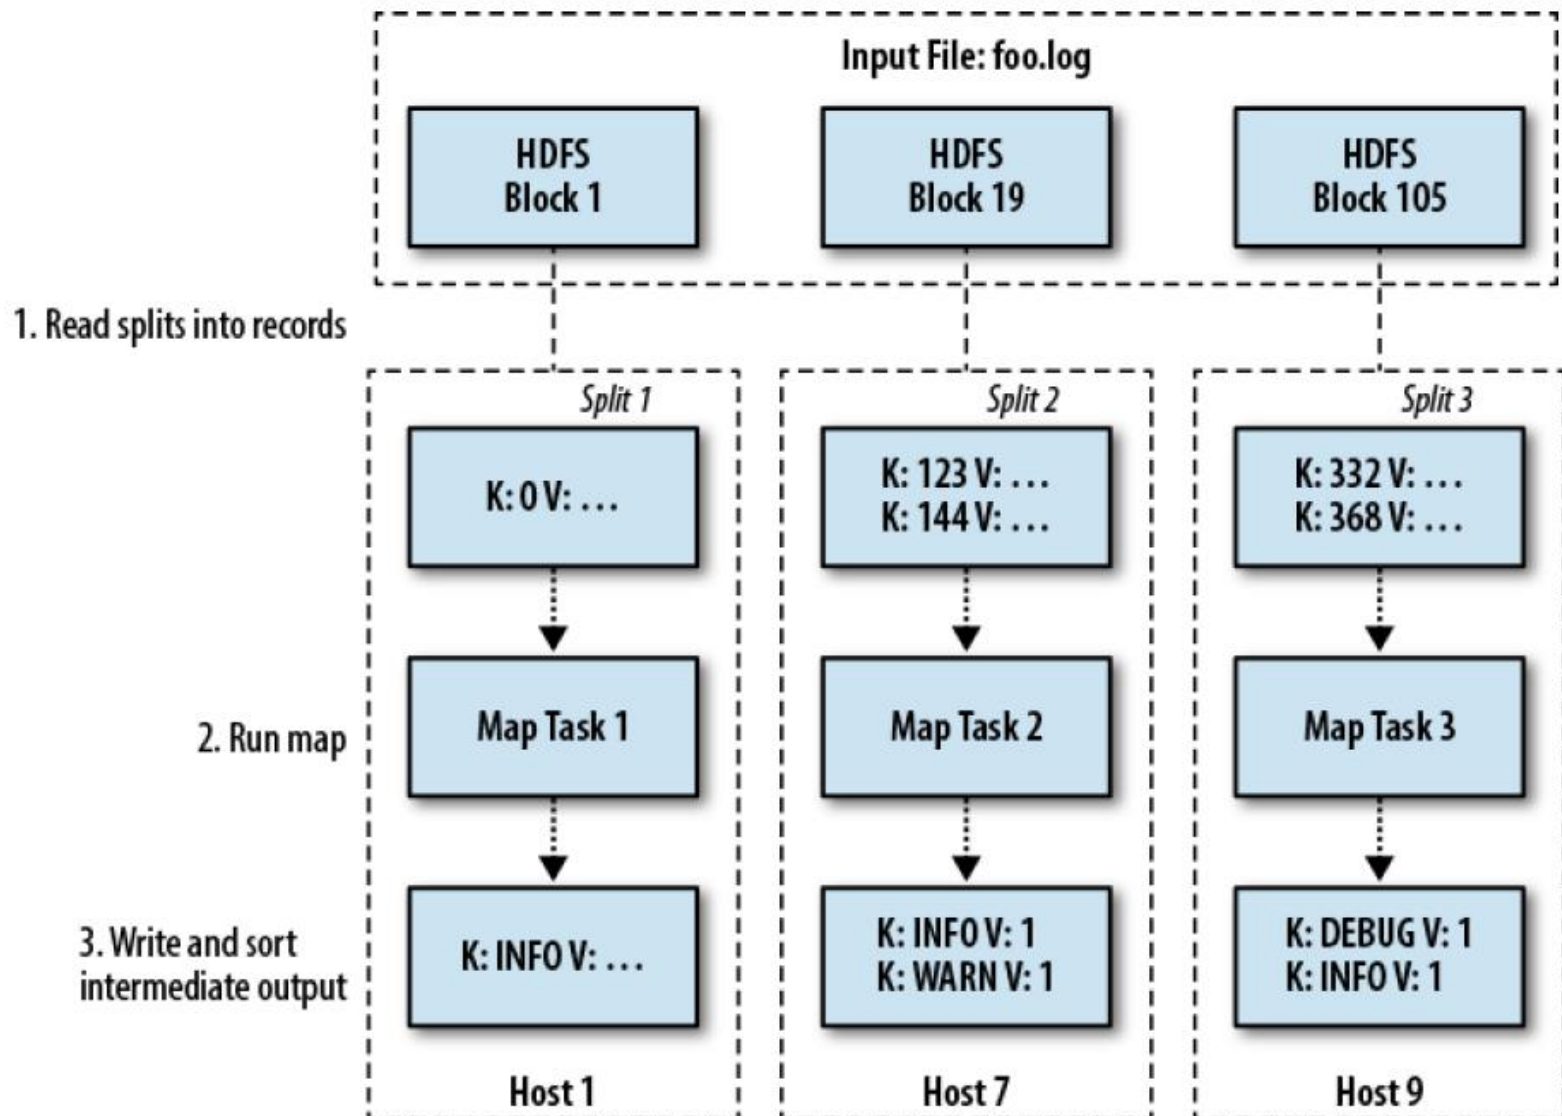
# MapReduce Example

- Word Count

# Word Count - Mapper

- 將輸入的文字檔案切成split
  - 每個mapper負責一個split
  - 由InputFormat決定有多少個split
- Mapper處理split中的每一筆record
  - 由RecordReader定義一筆key/value record
- 將每一筆record內的字輸出 (字, 1)
  - 真正map()所執行的工作





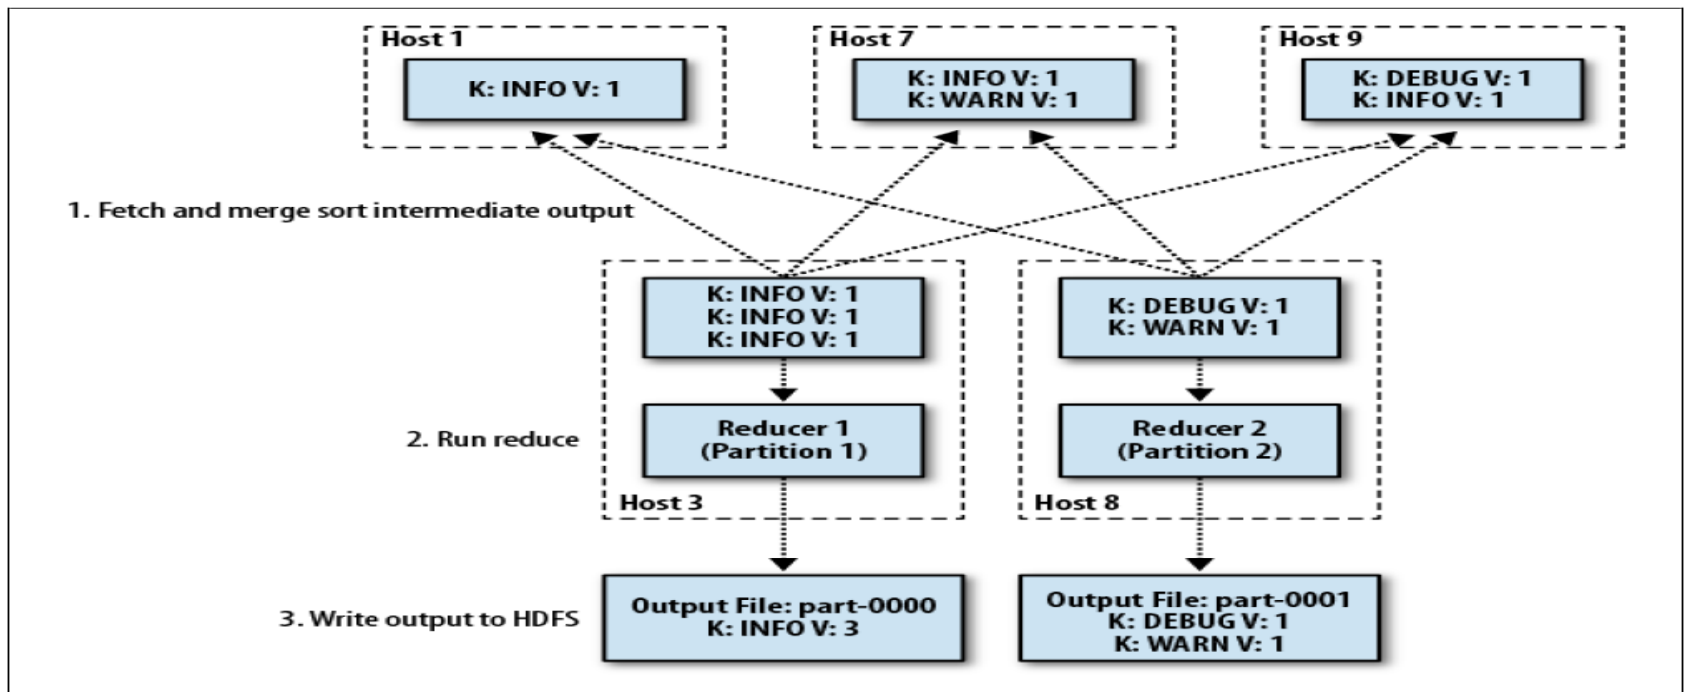
# Word Count – Shuffle & Sort

- Black box
  - 開發人員不用煩腦，framework會自行處理
- 在給Reducer之前完成
- 保證Reducer得到的資訊有下列三個特性
  - 可有多個Reducer
  - 同一個Reducer有可能處理多個Key值
  - 若Reducer看到某個Key1，會看到相對應的所有value
    - 給定Key1，所有Key1的值都會被同一個Reducer處理
    - Reducer 1收到 This這個字，會收到很多 1



# Word Count - Reducer

- Reducer收到許多 key與相對應的value list
  - Reducer1 收到 ( INFO, [1,1,1] )
  - Reducer 2 收到 (DEBUG, [1]), (WARN, [1])
  - Reducer 對每個字的出現次數做加總



# 正規描述

- Mapper
  - $(k1, v1) \rightarrow \text{list}(k2, v2)$
  - $(0, \text{"This is a book book"}) \rightarrow$   
 $(\text{"This"}, 1), (\text{"is"}, 1), (\text{"a"}, 1), (\text{"book"}, 1), (\text{"book"}, 1)$
- Reducer
  - $(k2, \text{list}(v2)) \rightarrow (k3, v3)$
  - $(\text{"This"}, [1]) \rightarrow (\text{"This"}, 1)$
  - $(\text{"is"}, [1]) \rightarrow (\text{"is"}, 1)$
  - $(\text{"a"}, [1]) \rightarrow (\text{"a"}, 1)$
  - $(\text{"book"}, [1, 1]) \rightarrow (\text{"book"}, 2)$

# Word Count – Pseudo code

```
void Map (key, value){  
    for each word x in value:  
        output.collect(x, 1);  
}
```

```
void Reduce (keyword, <list of value>){  
    for each x in <list of value>:  
        sum+=x;  
    final_output.collect(keyword, sum);  
}
```

split

map

shuffle

Partition  
& sort

grouping

reduce

This is a book  
That is a desk

This 1  
is 1  
a 1  
book 1  
That 1  
is 1  
a 1  
desk 1

I have a book

I 1  
have 1  
a 1  
book 1

I have a desk

I 1  
have 1  
a 1  
desk 1

a 1  
a 1  
a 1  
a 1  
book 1  
book 1  
desk 1  
desk 1  
have 1  
have 1

is 1  
is 1  
I 1  
I 1  
That 1  
This 1

a [1,1,1,1]  
book [1,1]  
desk [1,1]  
have [1,1]

is [1,1]  
I [1,1]  
That [1]  
This [1]


a 4  
book 2  
desk 2  
have 2

is 2  
I 2  
That 1  
This 1

# Labs 0: IntelliJ IDEA Setup



- Install IntelliJ IDEA Community Version
- Download labs code
  - <https://github.com/ogre0403/hust-workshop-2016>
- Import labs project into IntelliJ IDEA

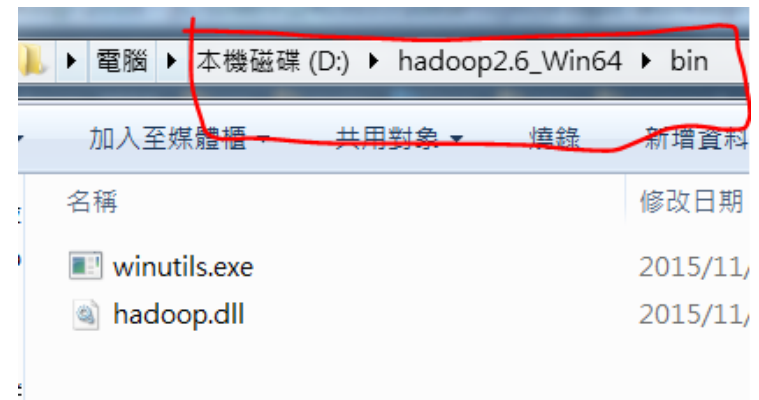



# Labs 0: IntelliJ IDEA Setup



- Windows 64
  - Setup winutils.exe/hadoop.dll

```
WordCountDriver
2016-10-25 15:07:12,752 [main] ERROR org.apache.hadoop.util.Shell - Failed to locate the winutils binary in the hadoop binary path
java.io.IOException: Could not locate executable null\bin\winutils.exe in the Hadoop binaries.
    at org.apache.hadoop.util.Shell.getQualifiedBinPath(Shell.java:355)
    at org.apache.hadoop.util.Shell.getWinUtilsPath(Shell.java:370)
    at org.apache.hadoop.util.Shell.<clinit>(Shell.java:363)
    at org.apache.hadoop.util.GenericOptionsParser.preProcessForWindows(GenericOptionsParser.java:438)
    at org.apache.hadoop.util.GenericOptionsParser.parseGeneralOptions(GenericOptionsParser.java:484)
    at org.apache.hadoop.util.GenericOptionsParser.<init>(GenericOptionsParser.java:170)
    at org.apache.hadoop.util.GenericOptionsParser.<init>(GenericOptionsParser.java:153)
    at ans.mapreduce.wordcount.WordCountDriver.main(WordCountDriver.java:20) <5 internal calls>
Usage: <input> <output>
```



# Labs 1: Word Count(MapReduce)



- WordCountMapper()
  - Use StringTokenizer.nextToken() to get word, then set the value by Text.set()
- WordCountReducer
  - Iterate all elements in values, and sum up all IntWritable objects


# MapReduce Example

- K-means



# K-means clustering

- 隨機選取資料組中的 $k$ 筆資料當作初始群中心 $u_1 \sim u_k$
- 計算每個資料 $x_i$ 對應到最短距離的群中心 (固定  $u_i$  求解所屬群  $S_i$ )
- 利用目前得到的分類重新計算群中心 (固定  $S_i$  求解群中心  $u_i$ )
- 重複step 2,3直到收斂 (達到最大疊代次數 or 群心中移動距離很小)



# 集中式版本程式

```
// Add in new data, one at a time, recalculating centroids with each new one.
while(!finish) {
    //Clear cluster state
    clearClusters();

    List lastCentroids = getCentroids();

    //Assign points to the closer cluster
    assignCluster();

    //Calculate new centroids.
    calculateCentroids();

    iteration++;

    List currentCentroids = getCentroids();

    //Calculates total distance between new and old Centroids
    double distance = 0;
    for(int i = 0; i < lastCentroids.size(); i++) {
        distance += Point.distance(lastCentroids.get(i),currentCentroids.get(i));
    }
    System.out.println("#####");
    System.out.println("Iteration: " + iteration);
    System.out.println("Centroid distances: " + distance);
    plotClusters();

    if(distance == 0) {
        finish = true;
    }
}
```

## Map

輸入為<目前的中心，point>

求point到每個中心的距離

輸出為<所屬的中心，point>

Read Distributed cache

C1 : (x1,y1)

C2 : (x2,y2)

C3 : (x3,y3)

Key	value
<hr/>	
C0	V1(1,2)
C0	V2(7,4)
C0	V3(16,3)
C0	V4(-1,-23)



mapper



Key	value
<hr/>	
C2	V1(1,2)

Key	value
<hr/>	
C2	V2(7,4)

Key	value
<hr/>	
C1	V3(16,3)

Key	value
<hr/>	
C3	V4(-1,-23)

## Reducer

輸入為<中心，屬於該中心的所有point>

對所有的point計算出新的中心

輸出<新的中心，point>做為下一次疊代

Key	value
-----	
C1	V3(16,3)

Key	value
-----	
C2	V1(1,2)
C2	V2(7,4)

Key	value
-----	
C3	V4(-1,-23)



reducer



Key	value
-----	
C1	V3(16,3)
C2	V1(1,2)
C2	V2(7,4)
C3	V4(-1,-23)

Update Distributed cache

C1 : (x'1,y'1)

C2 : (x'2,y'2)

C3 : (x'3,y'3)

# Labs 2: K-means (MapReduce)




- KMeansMapper
  - Use `DistanceMeasurer.measureDistance()` to calculate distance between vector and `ClusterCenter`
  - Find the nearest `ClusterCenter` and the shortest distance
- KMeansReducer
  - Sum up all Vector value. (Each digital is stored in `source[]`)  
Save result in `resultVector[]`.
  - Calculate mean of each digital in `resultVector[]`

# MapReduce Example

## - Page Rank

# PageRank

- 評估網頁重要程度的指標



$$PR(A) = \frac{PR(B)}{L(B)} + \frac{PR(C)}{L(C)} + \frac{PR(D)}{L(D)}$$
$$= \sum_{p_j} \frac{\text{PageRank}(p_j)}{L(p_j)}$$

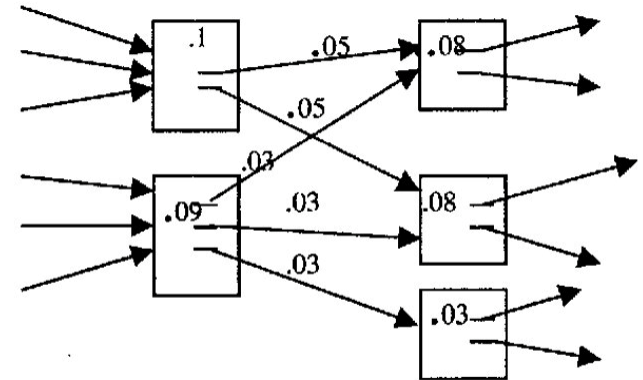



图 1 链接结构中的部分网页及其 PageRank 值



Page link matrix

= adjacency matrix

$M[i][j] = 1$  表示由  $i$  到  $j$  有一個邊

	A	B	C	D
A	0	1	1	1
B	1	0	0	1
C	1	0	0	0
D	0	1	1	0

Page link probability matrix

$M[i][j] = p$  表示由  $i$  到  $j$  的機率為  $p$

	A	B	C	D
A	0	1/3	1/3	1/3
B	1/2	0	0	1/2
C	1	0	0	
D	0	1/2	1/2	0

Transport Page link probability matrix

$M[i][j] = p$  表示由  $j$  到  $i$  的機率為  $p$   
 = 前一頁中的  $1 / L_j$

	A	B	C	D
A	0	1/2	1	0
B	1/3	0	0	1/2
C	1/3	0	0	1/2
D	1/3	1/2	0	0



$$\begin{aligned} \text{PR}(A) = & P(B \rightarrow A) * \text{PR}(B) \\ & + P(C \rightarrow A) * \text{PR}(C) \\ & + P(D \rightarrow A) * \text{PR}(D) \end{aligned}$$

$$\begin{aligned} \text{PR}(B) = & P(A \rightarrow B) * \text{PR}(A) \\ & + P(C \rightarrow B) * \text{PR}(C) \\ & + P(D \rightarrow B) * \text{PR}(D) \end{aligned}$$

$$\begin{aligned} \text{PR}(C) = & P(A \rightarrow C) * \text{PR}(A) \\ & + P(B \rightarrow C) * \text{PR}(B) \\ & + P(D \rightarrow C) * \text{PR}(D) \end{aligned}$$

$$\begin{aligned} \text{PR}(D) = & P(A \rightarrow D) * \text{PR}(A) \\ & + P(B \rightarrow D) * \text{PR}(B) \\ & + P(C \rightarrow D) * \text{PR}(C) \end{aligned}$$

$$\sum_{p_j} \frac{\text{PageRank}(p_j)}{L(p_j)}$$

Iteration 1

P	A	B	C	D
A	0	1/2	1	0
B	1/3	0	0	1/2
C	1/3	0	0	1/2
D	1/3	1/2	0	0

X

	PR
A	1/4
B	1/4
C	1/4
D	1/4

=

	PR
A	9/24
B	5/24
C	5/24
D	5/24

Iteration 2

P	A	B	C	D
A	0	1/2	1	0
B	1/3	0	0	1/2
C	1/3	0	0	1/2
D	1/3	1/2	0	0

X

	PR
A	9/24
B	5/24
C	5/24
D	5/24

=

	PR
A	15/48
B	11/48
C	11/48
D	11/48


Iteration N

P	A	B	C	D
A	0	1/2	1	0
B	1/3	0	0	1/2
C	1/3	0	0	1/2
D	1/3	1/2	0	0

...

	PR
A	3/9
B	2/9
C	2/9
D	2/9


# Reduce Pseudo Code



- Input:
  - Key:  $\langle p_0 \rangle$
  - Value:  $[\langle p_1, p_2, \dots, p_n \rangle, \langle PR, L \rangle, \langle PR, L \rangle \dots]$
- Output
  - Key:  $\langle p_0 \text{ new\_PR} \rangle$
  - Value:  $\langle p_1, p_2, \dots, p_n \rangle$

# Map Pseudo Code

- Input:
  - Key:  $\langle p0, PR \rangle$
  - Value:  $\langle p1, p2, \dots, pn \rangle$
- Output:
  - Type 1
    - Key:  $\langle p1 \rangle (\langle p2 \rangle, \langle p3 \rangle, \dots)$
    - Value:  $\langle PR L \rangle$
  - Type 2
    - Key:  $\langle p0 \rangle$
    - Value:  $\langle p1, p2, \dots, pn \rangle$



A, 0.25	B,C,D
---------	-------

B, 0.25	A, D
---------	------

C, 0.25	A
---------	---

D, 0.25	B,C
---------	-----

B	0.25, 3
C	0.25, 3
D	0.25, 3
A	B,C,D

A	0.25, 2
D	0.25, 2
B	A, D

A	0.25, 1
C	A

B	0.25, 2
C	0.25, 2
D	B, C

A	0.25, 2
	0.25, 1
	B,C,D

B	0.25, 3
	0.25, 2
	A, D

C	0.25, 3
	0.25, 2
	A

D	0.25, 3
	0.25, 2
	D, C

A, 0.375	B,C,D
----------	-------

B, 0.208	A, D
----------	------

C, 0.208	A
----------	---

D, 0.208	B,C
----------	-----

# Labs 3: Page Rank (MapReduce)





- RankCalculateMapper
  - For each linked to page, store (page,  $\frac{\text{thisPagesRank}}{\text{TotalLinksNumber}}$  )
- RankCalculateReducer
  - Calculate fraction pagerank contributed from linked page.
  - Sum up all contributed pagerank.

# Hadoop v.s. Spark

# Why Spark

- Compare with Hadoop ecosystem
  - More efficient execution
  - More unified program abstraction
  - More flexible program operation




transformation
<code>map(func)</code>
<code>filter(func)</code>
<code>flatMap(func)</code>
<code>sample(withReplacement, fraction, seed)</code>
<code>union(otherDataset)</code>
<code>distinct([numTasks])</code>
<code>groupByKey([numTasks])</code>
<code>reduceByKey(func, [numTasks])</code>
<code>sortByKey([ascending], [numTasks])</code>
<code>join(otherDataset, [numTasks])</code>
<code>cogroup(otherDataset, [numTasks])</code>
<code>cartesian(otherDataset)</code>

# RDD Concept



# RDD Essentials

- Resilient distributed dataset
- Each RDD is split into multiple **partitions**
  - Partitions may exist on different machines
- immutable distributed collection of objects
  - Transform creates new RDD
  - Coarse-grained transformation
- Spark keeps track lineage graph
  - Fast recovery from failure
- Lazy Evaluation
  - Until action is called



# RDD operations

## Transformations

- Create a new dataset from and existing one.
- Lazy in nature. They are executed only when some action is performed.
- Example :
  - Map(func)
  - Filter(func)
  - Distinct()


## Actions

- Returns to the driver program a value or exports data to a storage system after performing a computation.
- Example:
  - Count()
  - Reduce(func)
  - Collect
  - Take()

## Persistence

- For caching datasets in-memory for future operations.
- Option to store on disk or RAM or mixed (Storage Level).
- Example:
  - Persist()
  - Cache()


# Create RDD



```
// base RDD  
val lines = sc.textFile("hdfs://...")
```

```
JavaRDD<String> lines = sc.textFile("/path/to/README.md");
```


# RDD Transformation



```
// transformed RDDs  
val errors = lines.filter(_.startsWith("ERROR"))  
val messages = errors.map(_.split("\t")).map(r => r(1))  
messages.cache()
```

```
JavaRDD<String> errorsRDD = inputRDD.filter(  
    new Function<String, Boolean>() {  
        public Boolean call(String x) { return x.contains("error"); }  
    }  
));
```

# RDD Action




```
// action 1  
messages.filter(_.contains("mysql")).count()
```

```
System.out.println("Input had " + badLinesRDD.count() + " concerning lines")  
System.out.println("Here are 10 examples:")  
for (String line: badLinesRDD.take(10)) {  
    System.out.println(line);  
}
```

# Lineage Graph

- Think of each RDD as consisting of instructions on how to compute the data through transformations.

```
errorsRDD = inputRDD.filter(lambda x: "error" in x)
warningsRDD = inputRDD.filter(lambda x: "warning" in x)
badLinesRDD = errorsRDD.union(warningsRDD)
```



# RDD Operations

# RDD Type

- Basic RDD[T]
  - considers each data item as a single value
- Convert to other RDD Type
- PairRDD[K,V]
  - each data item containing key/value pairs.

<http://blog.csdn.net/pelick/article/details/44922619>

<http://homepage.cs.latrobe.edu.au/zhe/ZhenHeSparkRDDAPIExamples.html>



# Basic RDD Transformation

- `RDD.map(func)`

*Table 3-2. Basic RDD transformations on an RDD containing {1, 2, 3, 3}*

Function name	Purpose	Example	Result
<code>map()</code>	Apply a function to each element in the RDD and return an RDD of the result.	<code>rdd.map(x =&gt; x + 1)</code>	<code>{2, 3, 4, 4}</code>

```
JavaRDD<Integer> rdd = sc.parallelize(Arrays.asList(1, 2, 3, 4));
JavaRDD<Integer> result = rdd.map(new Function<Integer, Integer>() {
    public Integer call(Integer x) { return x*x; }
});
System.out.println(StringUtils.join(result.collect(), ","));
```

# Basic RDD Transformation

- `RDD.filter(func)`

*Table 3-2. Basic RDD transformations on an RDD containing {1, 2, 3, 3}*

Function name	Purpose	Example	Result
<code>filter()</code>	Return an RDD consisting of only elements that pass the condition passed to <code>filter()</code> .	<code>rdd.filter(x =&gt; x != 1)</code>	<code>{2, 3, 3}</code>

```
RDD<String> errors = lines.filter(new Function<String, Boolean>() {  
    public Boolean call(String x) { return x.contains("error"); }  
});
```

# Basic RDD Transformation

- `RDD.flatMap(func)`

*Table 3-2. Basic RDD transformations on an RDD containing {1, 2, 3, 3}*

Function name	Purpose	Example	Result
<code>flatMap()</code>	Apply a function to each element in the RDD and return an RDD of the contents of the iterators returned. Often used to extract words.	<code>rdd.flatMap(x =&gt; x.to(3))</code>	{1, 2, 3, 2, 3, 3, 3}

```
JavaRDD<String> lines = sc.parallelize(Arrays.asList("hello world", "hi"));
JavaRDD<String> words = lines.flatMap(new FlatMapFunction<String, String>() {
    public Iterable<String> call(String line) {
        return Arrays.asList(line.split(" "));
    }
});
words.first(); // returns "hello"
```

# Basic RDD Action

- `RDD.reduce(func)`

*Table 3-4. Basic actions on an RDD containing {1, 2, 3, 3}*

Function name	Purpose	Example	Result
<code>reduce(func)</code>	Combine the elements of the RDD together in parallel (e.g., sum).	<code>rdd.reduce((x, y) =&gt; x + y)</code>	9

```
Integer sum = rdd.reduce(new Function2<Integer, Integer, Integer>() {  
    public Integer call(Integer x, Integer y) { return x + y; }  
});
```

- Java Basic RDD convert to
  - PairRDD : **explicitly** implement PairFunction()


Function name	Equivalent function*<A, B,...>	Usage
PairFunction<T, K, V>	Function<T, Tuple2<K, V>>	PairRDD<K, V> from a mapToPair
DoubleFunction<T>	Function<T, double>	DoubleRDD from map ToDouble

```
PairFunction<String, String, String> keyData =
    new PairFunction<String, String, String>() {
        public Tuple2<String, String> call(String x) {
            return new Tuple2(x.split(" ")[0], x);
        }
    };
JavaPairRDD<String, String> pairs = lines.mapToPair(keyData);
```

```
JavaDoubleRDD result = rdd.mapToDouble(
    new DoubleFunction<Integer>() {
        public double call(Integer x) {
            return (double) x * x;
        }
    });
System.out.println(result.mean());
```

# PairRDD Transformation

- PairRDD is also a RDD
  - RDD.filter(func)
  - RDD.map(func)
  - RDD.flatMap(func)
  - ...



```
Function<Tuple2<String, String>, Boolean> longWordFilter =  
    new Function<Tuple2<String, String>, Boolean>() {  
        public Boolean call(Tuple2<String, String> keyValue) {  
            return (keyValue._2().length() < 20);  
        }  
    };  
JavaPairRDD<String, String> result = pairs.filter(longWordFilter);
```

# PairRDD Transformation

- `RDD.mapValues(func)`

*Table 4-1. Transformations on one pair RDD (example: {(1, 2), (3, 4), (3, 6)})*

Function name	Purpose	Example	Result
<code>mapValues(func)</code>	Apply a function to each value of a pair RDD without changing the key.	<code>rdd.mapValues(x =&gt; x+1)</code>	{(1, 3), (3, 5), (3, 7)}

```
JavaPairRDD<Integer,Integer> result =  
    prdd.mapValues(new Function<Integer, Integer>() {  
        @Override  
        public Integer call(Integer v1) throws Exception {  
            return v1 +1;  
        }  
    });
```

# PairRDD Transformation

- `RDD.reduceByKey(func)`

*Table 4-1. Transformations on one pair RDD (example: {(1, 2), (3, 4), (3, 6)})*

Function name	Purpose	Example	Result
<code>reduceByKey(func)</code>	Combine values with the same key.	<code>rdd.reduceByKey( (x, y) =&gt; x + y)</code>	<code>{(1, 2), (3, 10)}</code>

```
JavaPairRDD<Integer,Integer> result =  
    prdd.reduceByKey(new Function2<Integer, Integer, Integer>() {  
        @Override  
        public Integer call(Integer v1, Integer v2) throws Exception {  
            return v1 + v2;  
        }  
    });
```

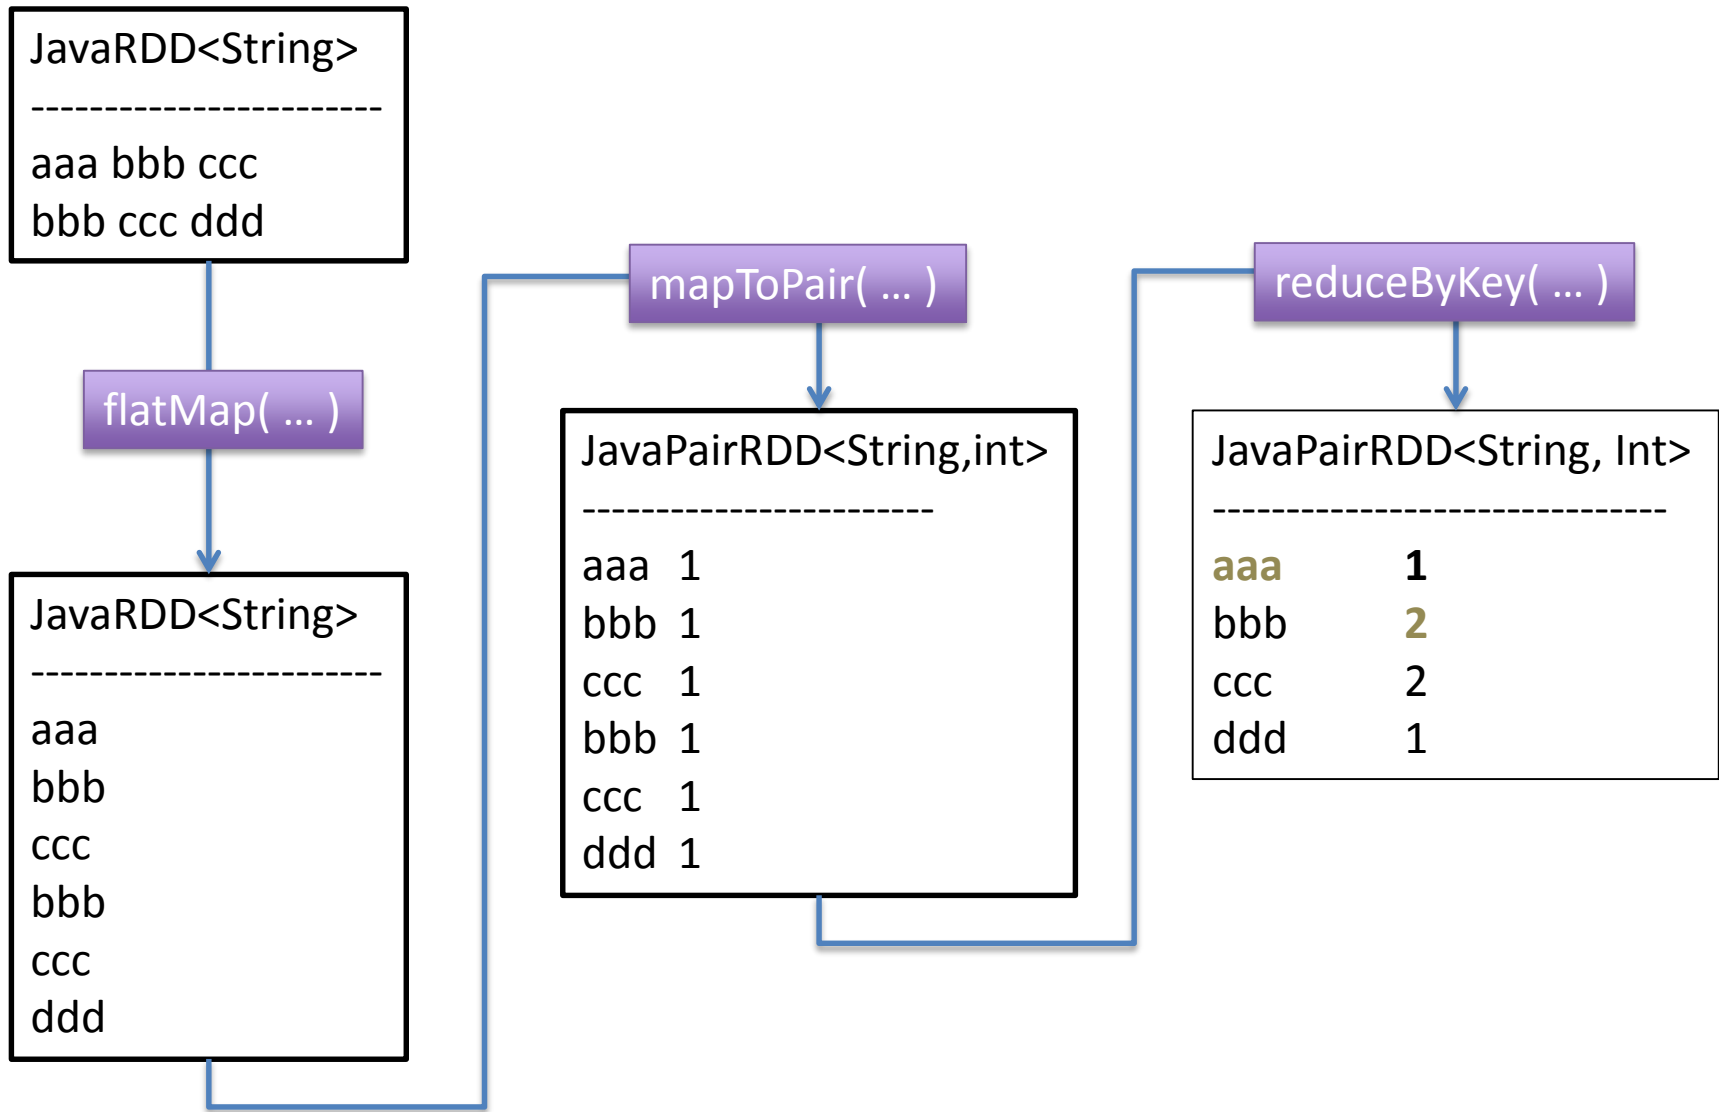


# Convert MR Job into Spark Job

- Map() in MR
  - flatmap() + map()/mapToPair()
- Reduce() in MR
  - reduceByKey()
  - groupByKey() + mapValue()

# Spark RDD example

- Word Count



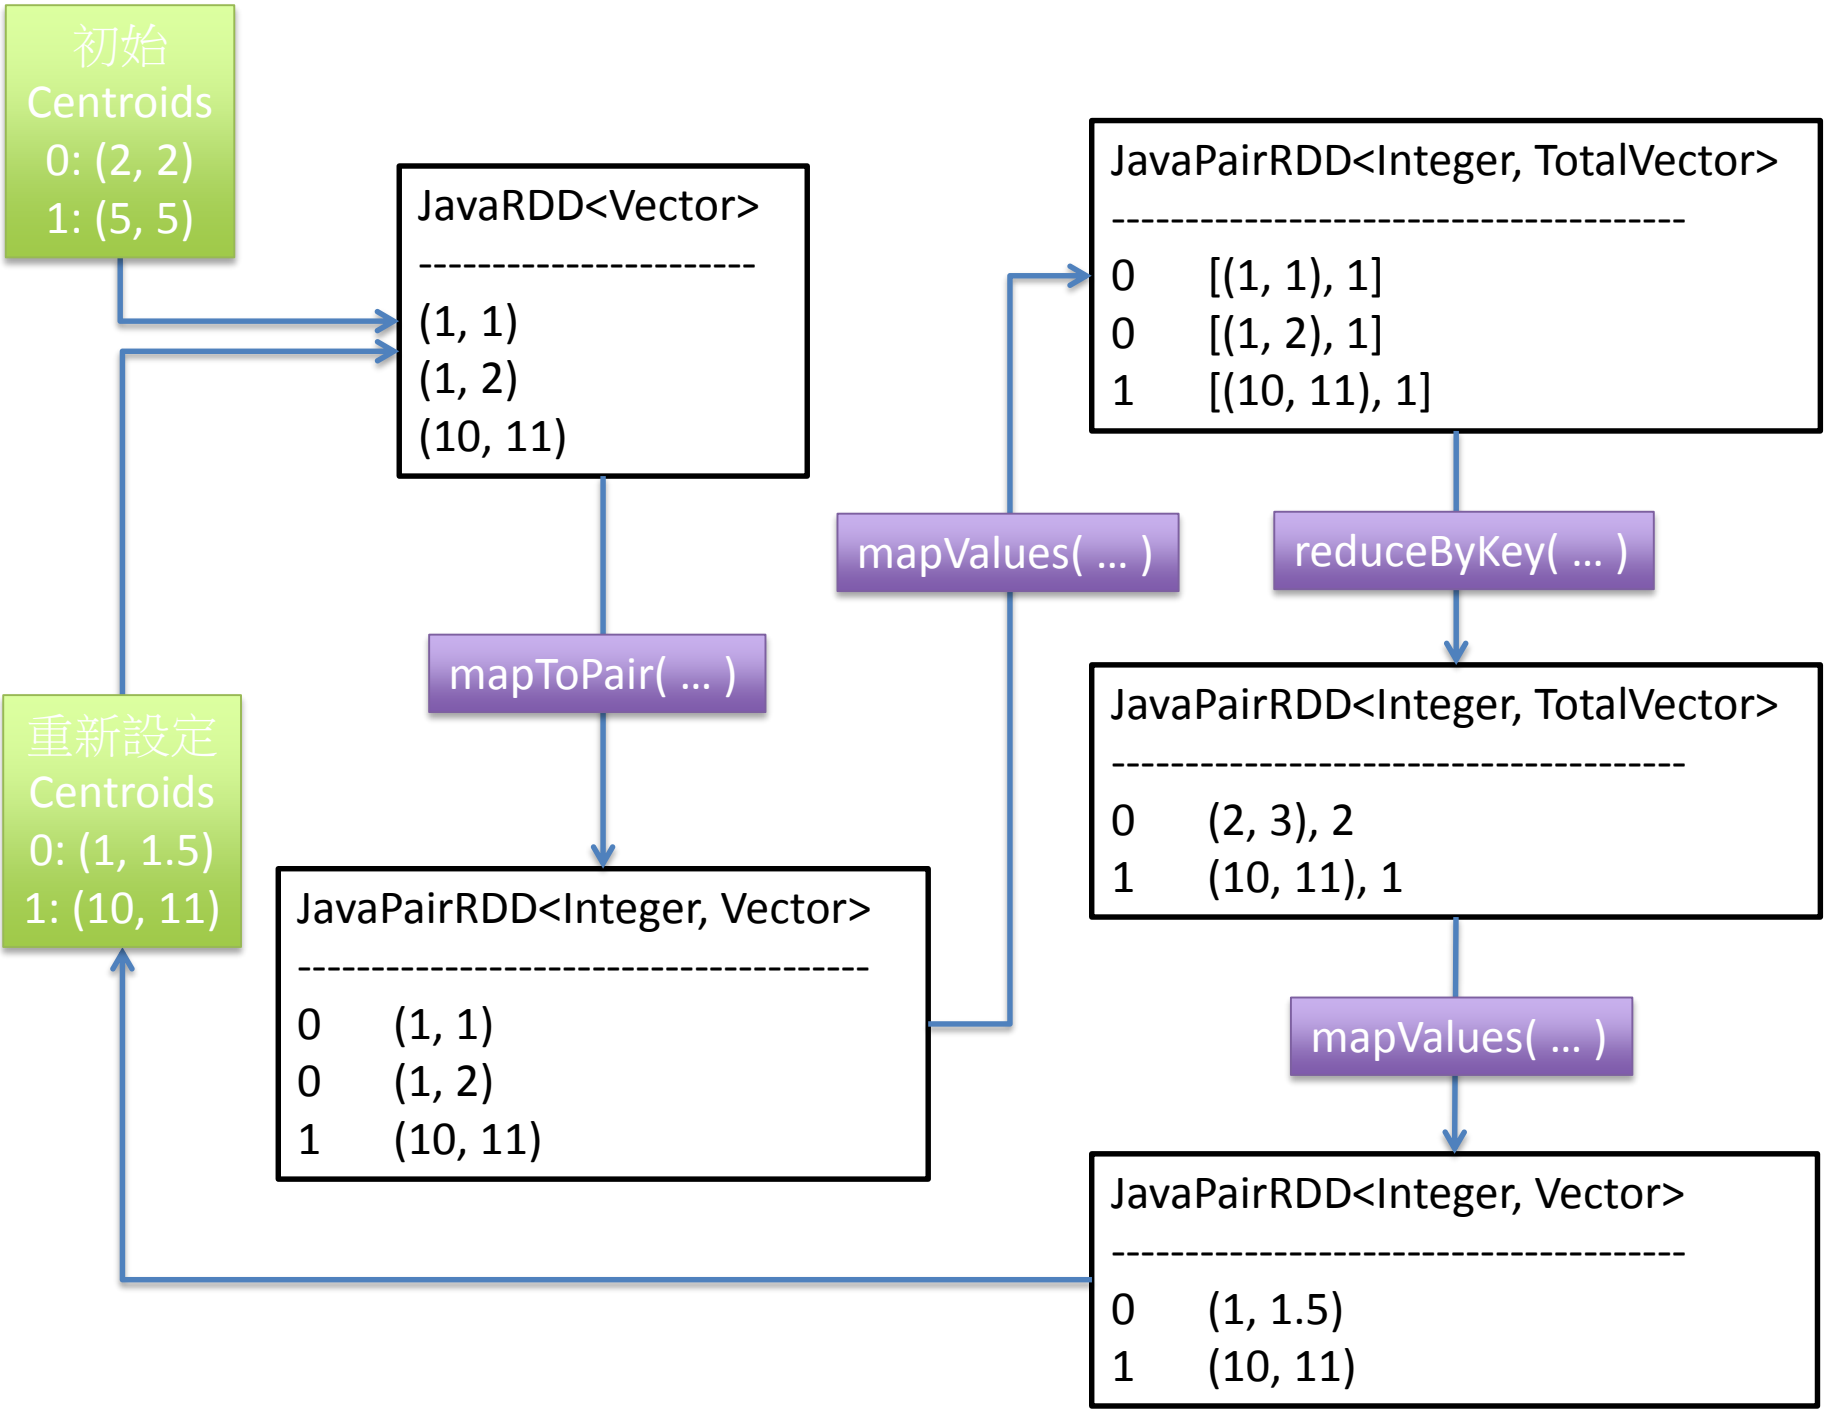
# Labs 4: Word Count(Spark)



- Use `flatMap()` to split each line into multiple words
  - Use `String.split()` to generate String array, then use `Arrays.asList()` to create String iterable
- Use `mapToPair()` to transform word into (word, one) pair
  - return (word, 1) tuple
- Use `reduceByKey()` to generate (word, count) pair
  - Create Function2 class
  - Implement Integer call(Integer v1, Integer v2)

# Spark RDD example

- K-means



# Labs 5: K-means (Spark)

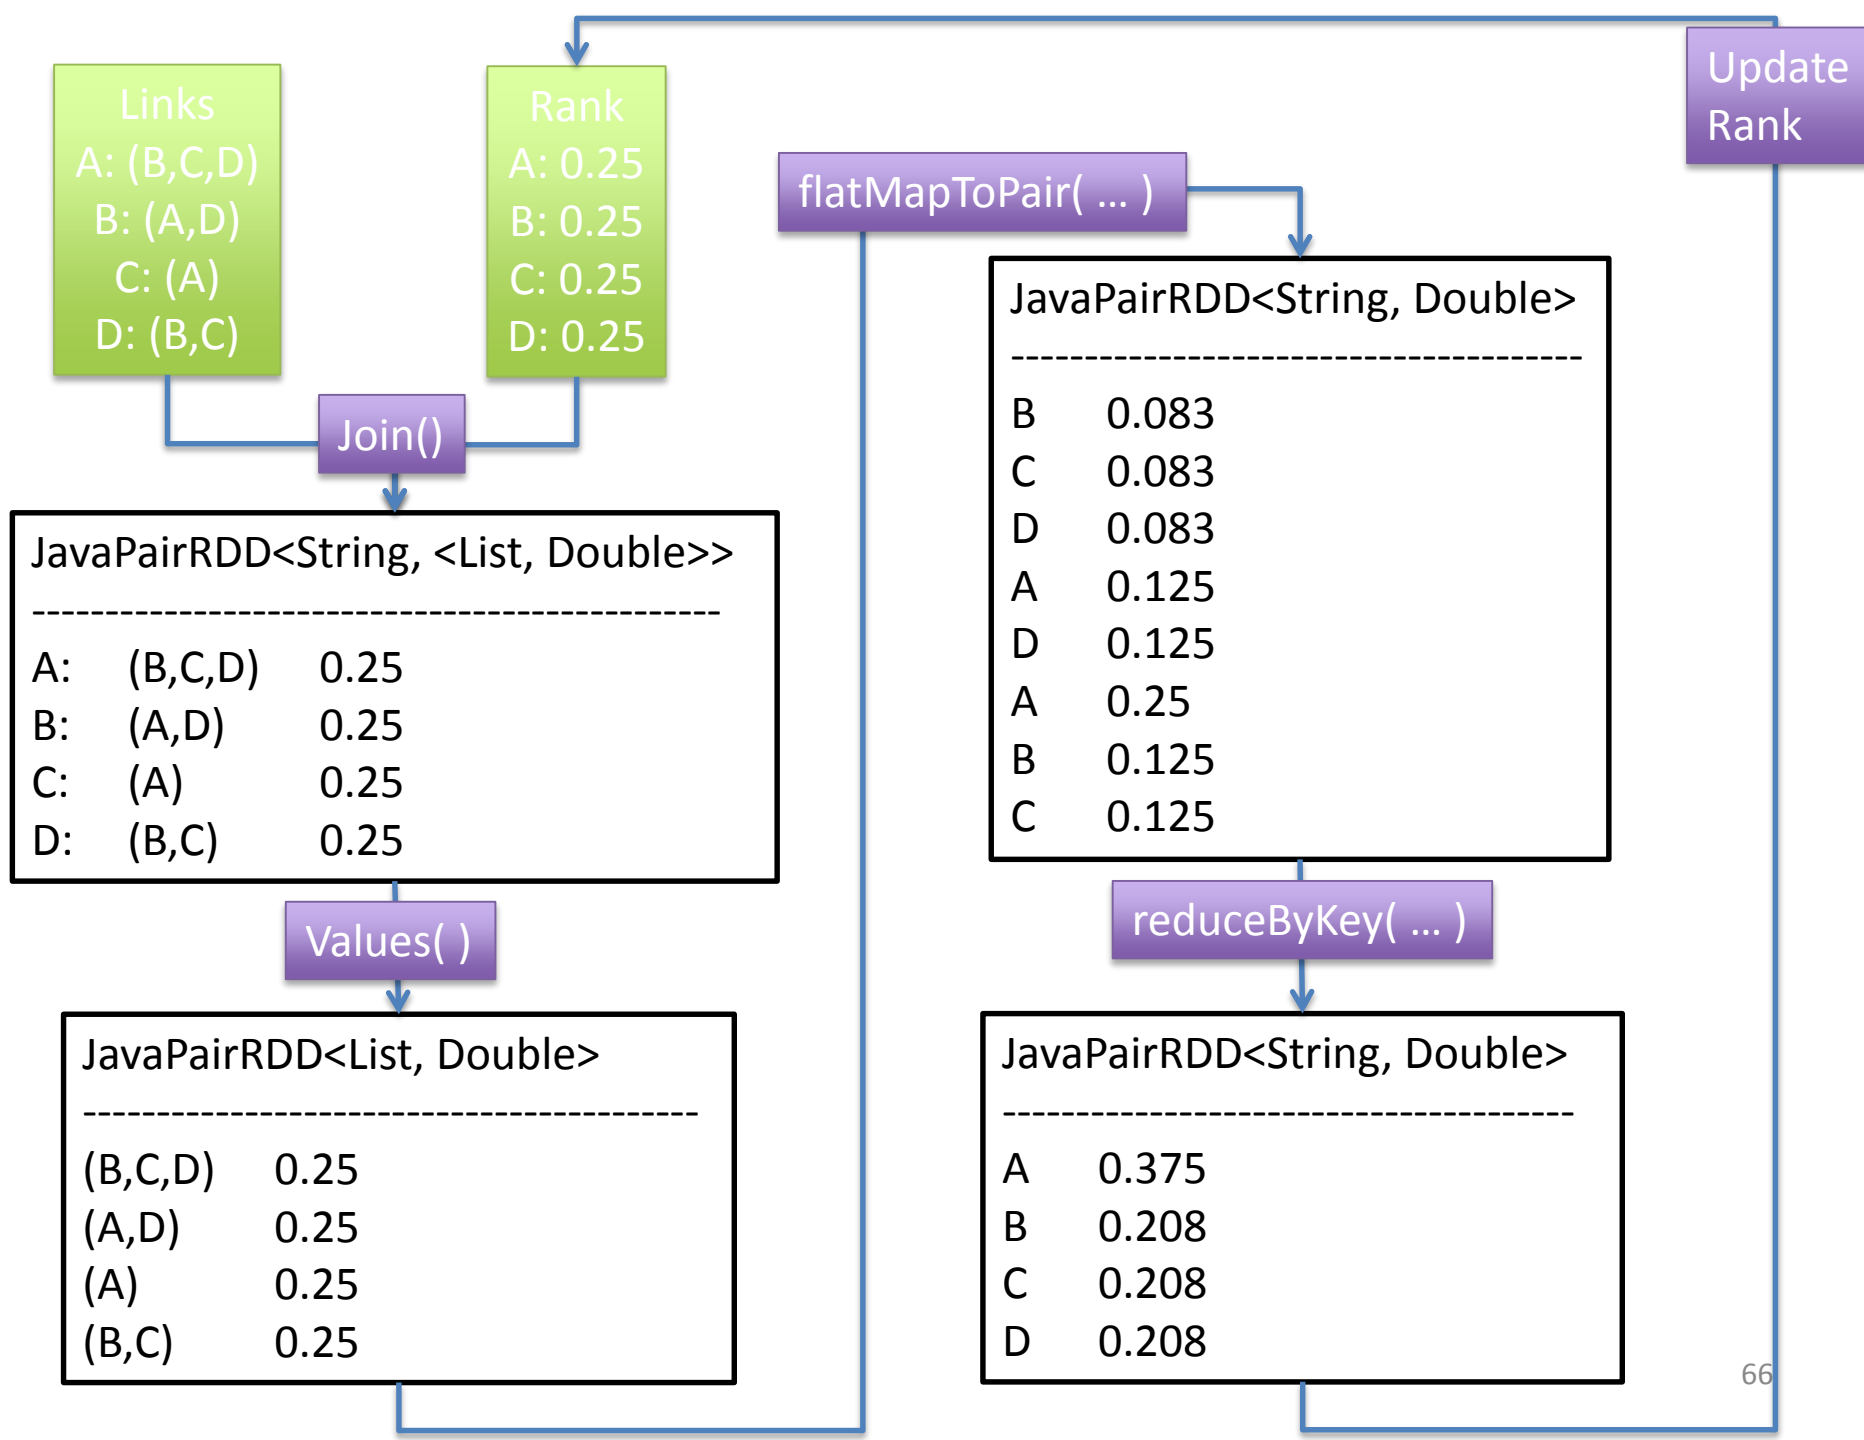


- 分群, 計算每個Vector離那一組中心最近
  - Use `closestPoint(v, centroids)` to find the ID of nearest center, then return (ID, Vector) tuple



# Spark RDD example

## - Page Rank



# Labs 6: Page Rank (Spark)



- For each URL, calculate the PR contributed from its neighbor, then add (url, PR) tuple into results list